

REDISCONF 2016

---

# BACKGROUND TASKS IN NODE.JS

@EVANTAHLE

**WARNING**

**MOST OF THE IDEAS IN THIS  
PRESENTATION ARE ACTUALLY VERY  
BAD IDEAS.**

**TRY THESE AT ~, NOT ON PRODUCTION**



## HI. I'M EVAN

- ▶ Director of Technology @ TaskRabbit
- ▶ Founder of ActionHero, node.js framework
- ▶ Node-Resque Author



@EVANTAHLER

[BLOG.EVANTAHLER.COM](http://BLOG.EVANTAHLER.COM)

# WHAT IS NODE.JS

- ▶ Server-side Framework, uses JS
- ▶ Async
- ▶ Fast

# WHAT IS REDIS

- ▶ In-memory database
- ▶ Structured data
- ▶ Fast



WHAT ARE WE GOING TO TALK ABOUT?

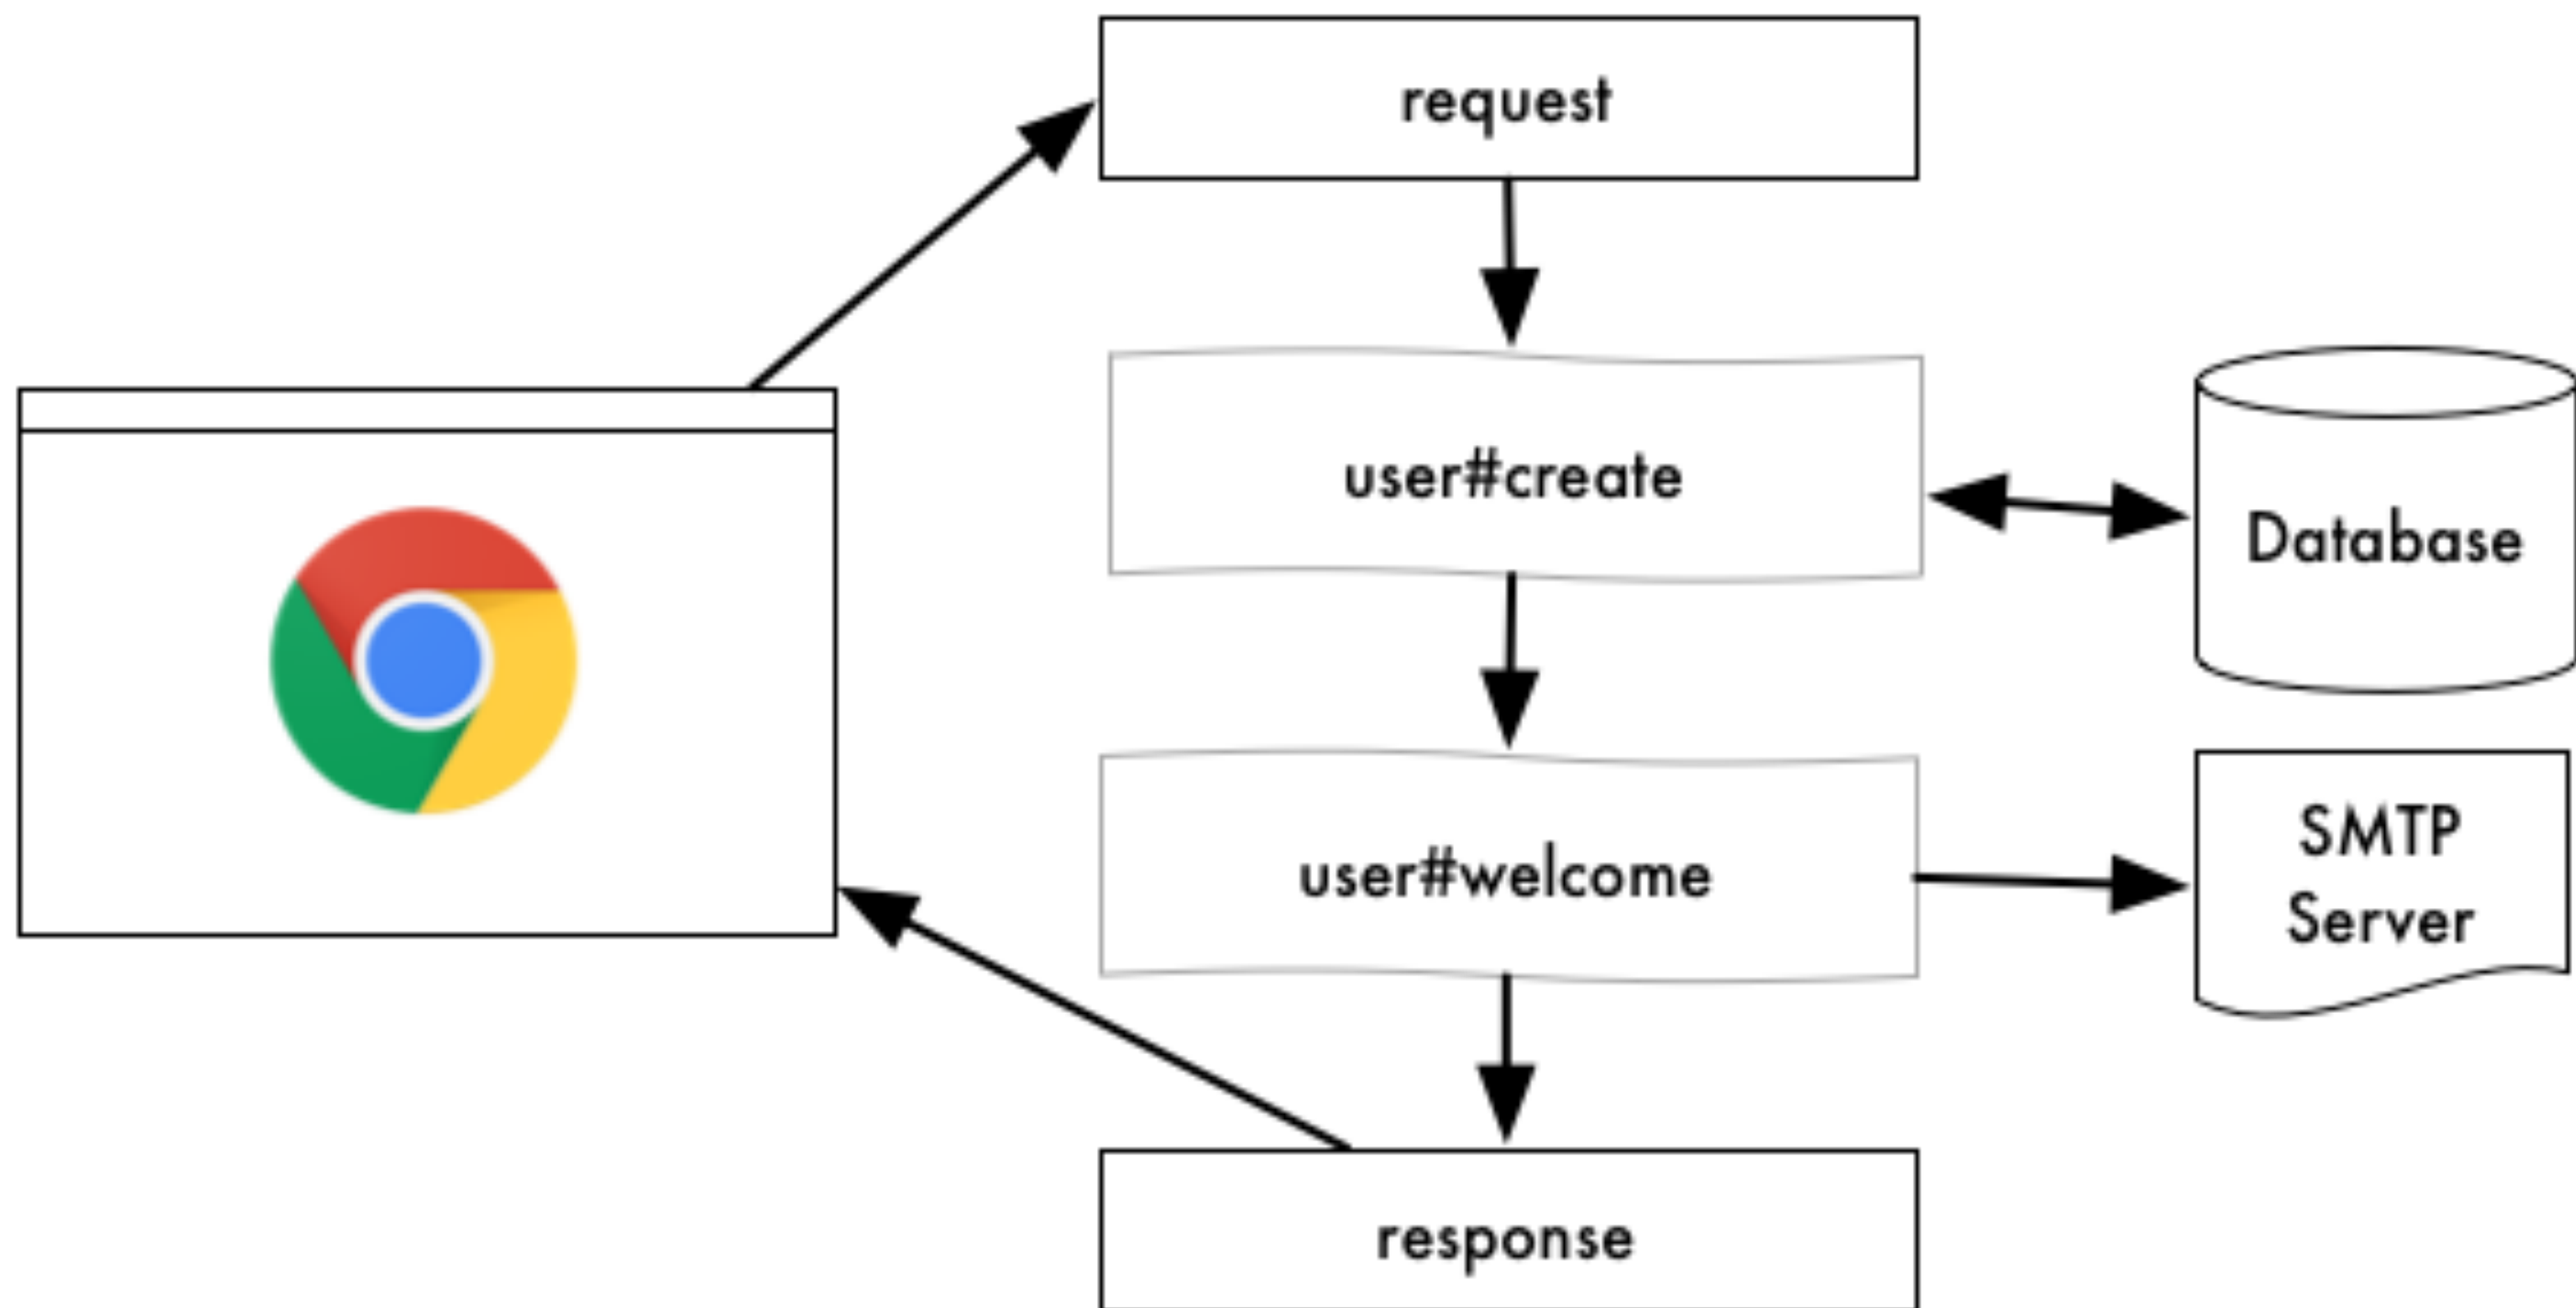
---

**EVERYTHING IS FASTER IN NODE...  
ESPECIALLY THE BAD IDEAS**

**Me (Evan)**

## WHAT ARE WE GOING TO TALK ABOUT?

---



# POSSIBLE TASK STRATEGIES:

FOREGROUND (IN-LINE)

PARALLEL (THREAD-ISH)

LOCAL MESSAGES (FORK-ISH)

REMOTE MESSAGES (\*MQ-ISH)

REMOTE QUEUE (REDIS + RESQUE)

IMMUTABLE EVENT BUS (KAFKA-ISH)



# 1) FOREGROUND TASKS

```
<?php
$to      = 'nobody@example.com';
$subject = 'the subject';
$message = 'hello';
$headers = 'From: webmaster@example.com' . "\r\n" .
    'Reply-To: webmaster@example.com' . "\r\n" .
    'X-Mailer: PHP/' . phpversion();

mail($to, $subject, $message, $headers);
?>
```

# SENDING EMAILS

```
var http          = require('http');
var nodemailer    = require('nodemailer');
var httpPort      = process.env.PORT || 8080;
var httpHost      = process.env.HOST || '127.0.0.1';

var transporter = nodemailer.createTransport({
  service: 'gmail',
  auth: {
    user: require('./.emailUsername'),
    pass: require('./.emailPassword')
  }
});
```

```
var sendEmail = function(req, callback){
  var urlParts = req.url.split('/');
  var email = {
    from: require('./.emailUsername'),
    to: decodeURI(urlParts[1]),
    subject: decodeURI(urlParts[2]),
    text: decodeURI(urlParts[3]),
  };
  transporter.sendMail(email, function(error, info){
    callback(error, email);
  });
};
```

# THE SERVER

```
var server = function(req, res){
  var start = Date.now();
  var responseCode = 200;
  var response      = {};

  sendEmail(req, function(error, email){
    response.email = email;

    if(error){
      console.log(error);
      responseCode = 500;
      response.error = error;
    }

    res.writeHead(responseCode, {'Content-Type': 'application/json'});
    res.end(JSON.stringify(response, null, 2));
    var delta = Date.now() - start;
    console.log('Sent an email to ' + email.to + ' in ' + delta + 'ms');
  });

};

http.createServer(server).listen(httpPort, httpHost);
```



Async and non-blocking!



**DEMO TIME**



# STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ The client still needs to wait for the message to send, but you won't block any other client's requests
  - ▶ Avg response time of ~2 seconds from my couch
- ▶ Why it is still a bad idea:
  - ▶ Slow for the client
  - ▶ Spending "web server" resources on sending email
  - ▶ Error / Timeout to the client for "partial success"
    - ▶ IE: Account created but email not sent
    - ▶ Confusing to the user, dangerous for the DB

## 2) PARALLEL TASKS

## IMPROVEMENT IDEAS

- ▶ In any other language this would be called “threading”
  - ▶ But if it were real threading, the client would still have to wait
  - ▶ I guess this might help you catch errors...
  - ▶ *\*note: do not get into a discussion about threads in node...*
- ▶ Lets get crazy:
  - ▶ Ignore the Callback




# IGNORE THE CALLBACK

```
var server = function(req, res){  
  var start = Date.now();  
  var responseCode = 200;  
  var response = {};  
  sendEmail(req);  
  res.writeHead(responseCode, {'Content-Type': 'application/json'});  
  res.end(JSON.stringify(response, null, 2));  
  console.log('Sent an email');  
};
```

```
http.createServer(server).listen(httpPort, httpHost);
```

```
var sendEmail = function(req, callback){  
  var urlParts = req.url.split('/');  
  var email = {  
    from: require('./emailUsername'),  
    to: decodeURI(urlParts[1]),  
    subject: decodeURI(urlParts[2]),  
    text: decodeURI(urlParts[3]),  
  };  
  transporter.sendMail(email, function(error, info){  
    if(typeof callback === 'function'){  
      callback(error, email);  
    }  
  });  
};
```





**DEMO TIME**



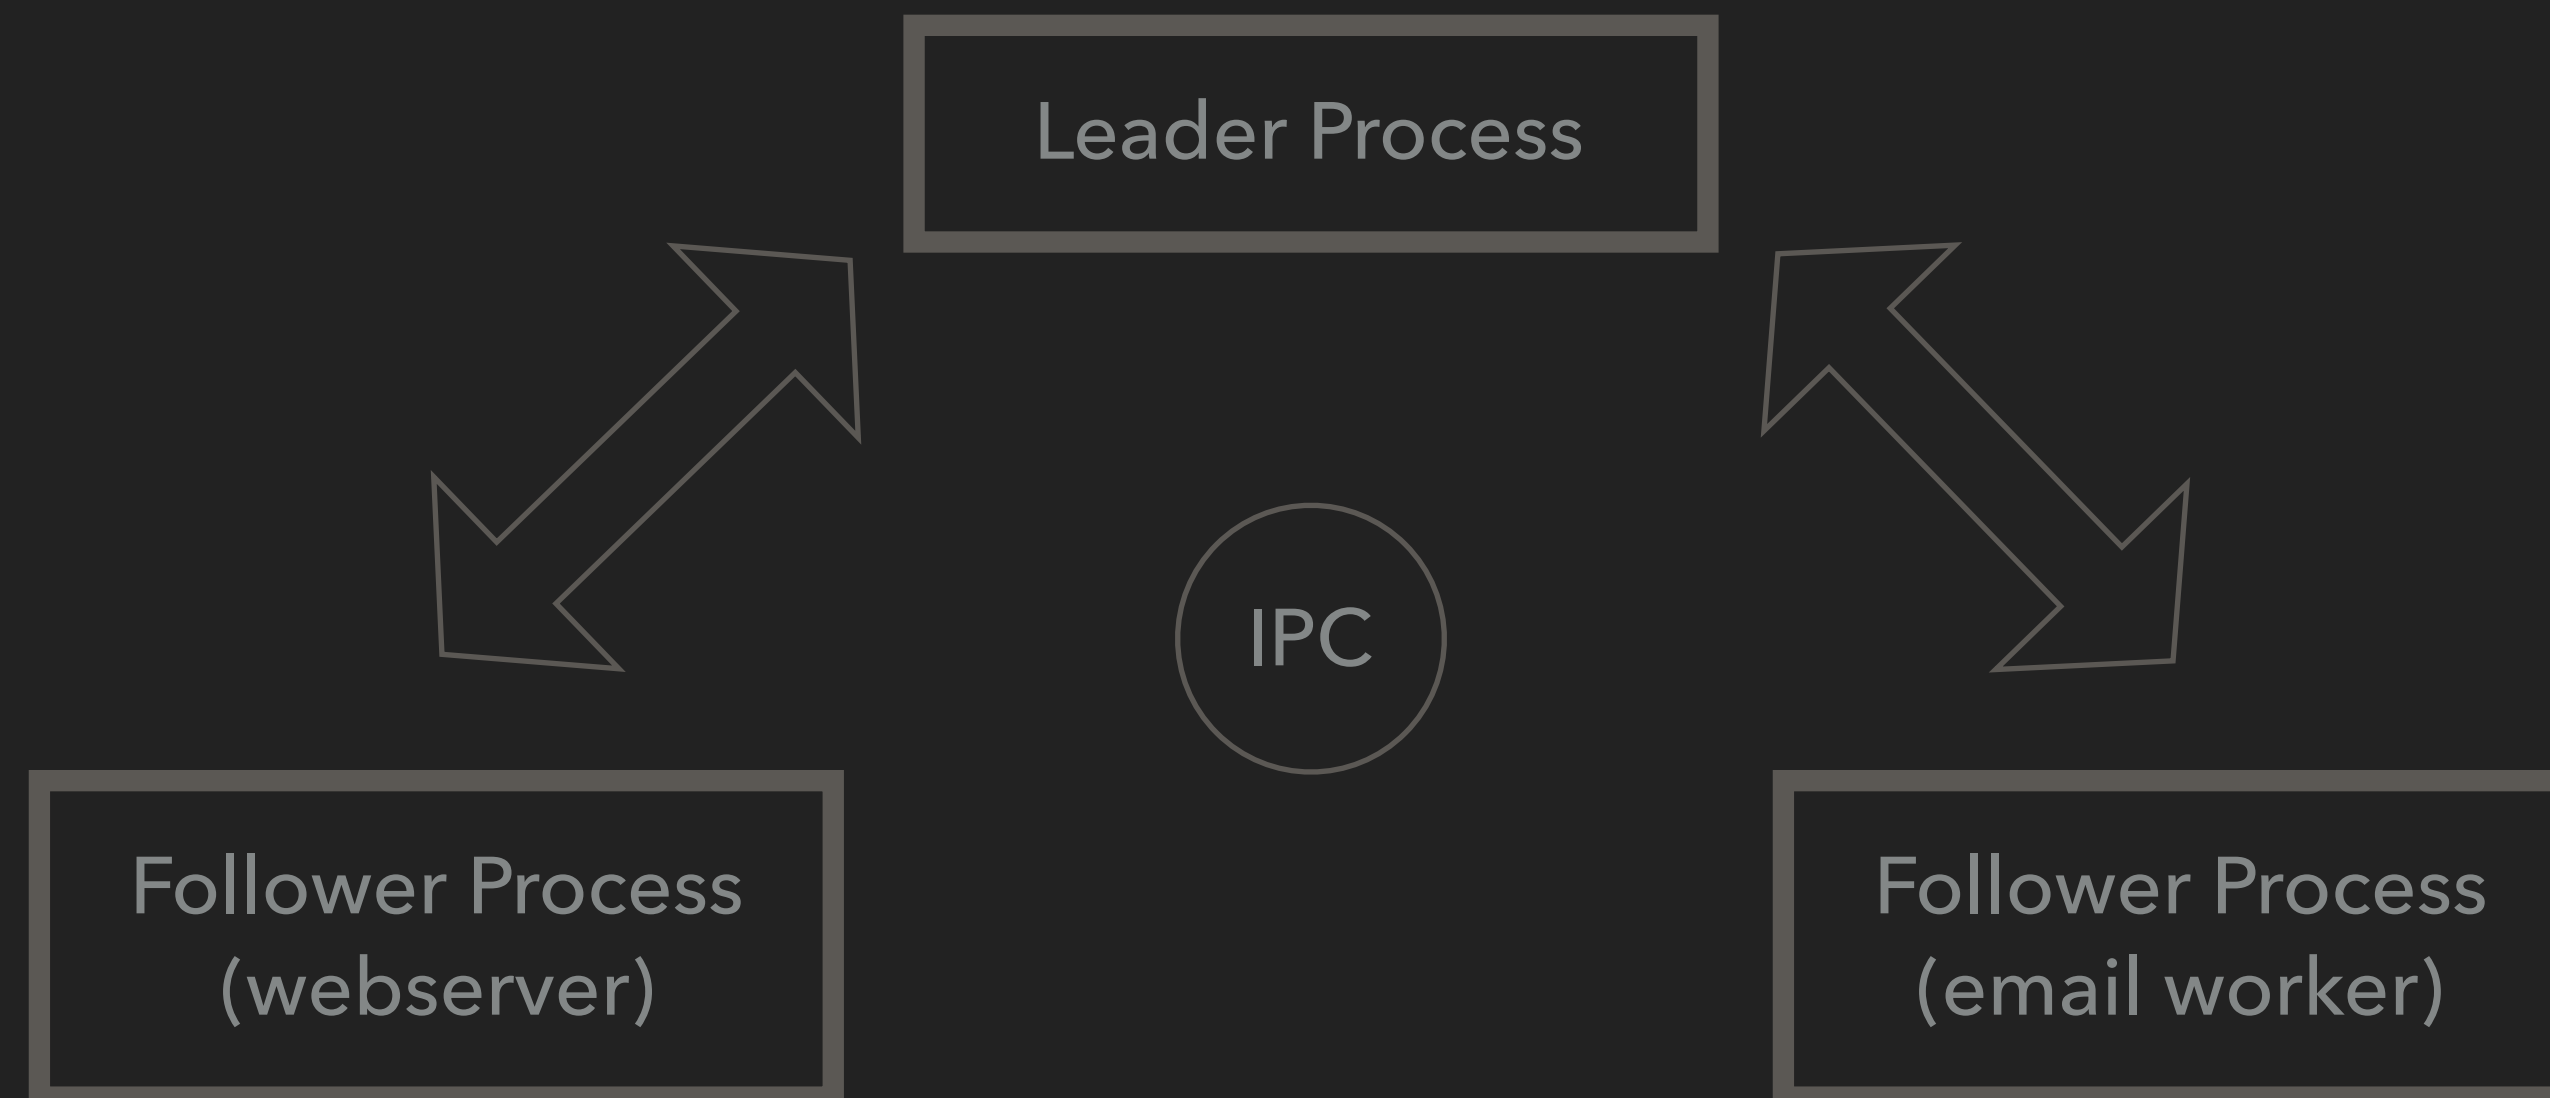
## STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ It's rare you can actually do this in a language... without threading or folks!
  - ▶ Crazy-wicked-fast.
- ▶ Why it is still a bad idea:
  - ▶ 0 callbacks, 0 data captured
  - ▶ I guess you could log errors?
  - ▶ But what would you do with that data?
  - ▶ The client has no idea what happened

# 3) LOCAL MESSAGES

*or: "The part of the talk where we grossly over-engineer some stuff"*

## IMPROVEMENT IDEAS



# CLUSTERING IN NODE.JS

## A SIMPLE TIMED MANAGER SCRIPT...

```
var cluster = require('cluster');

if(cluster.isMaster){
  doMasterStuff();
}else{
  if(process.env.ROLE === 'server'){ doServerStuff(); }
  if(process.env.ROLE === 'worker'){ doWorkerStuff(); }
}
```

```
var doMasterStuff = function(){
  log('master', 'started master');

  var masterLoop = function(){
    checkOnWebServer();
    checkOnEmailWorker();
  };

  var checkOnWebServer = function(){
    ...
  };

  var checkOnEmailWorker = function(){
    ...
  };

  setInterval(masterLoop, 1000);
};
```

# CLUSTERING IN NODE.JS

```
var doServerStuff = function(){
  var server = function(req, res){
    var urlParts = req.url.split('/');
    var email = {
      to: decodeURI(urlParts[1]),
      subject: decodeURI(urlParts[2]),
      text: decodeURI(urlParts[3]),
    };

    var response = {email: email};
    res.writeHead(200, {'Content-Type': 'application/json'});
    res.end(JSON.stringify(response, null, 2));

    process.send(email);
  };

  http.createServer(server).listen(httpPort, '127.0.0.1');
};
```



Interprocess Communication (IPC) with complex data-types

# CLUSTERING IN NODE.JS

```
var checkOnWebServer = function(){
  if(children.server === undefined){
    log('master', 'starting web server');
    children.server = cluster.fork({ROLE: 'server'});
    children.server.name = 'web server';
    children.server.on('online', function(){ log(children.server, 'ready on port ' + httpPort); });
    children.server.on('exit', function(){
      log(children.server, 'died :(');
      delete children.server;
    });
    children.server.on('message', function(message){
      log(children.server, 'got an email to send from the webserver: ' + JSON.stringify(message));
      children.worker.send(message);
    });
  }
};
```

...IT'S ALL JUST MESSAGE PASSING



# CLUSTERING IN NODE.JS

```
var checkOnEmailWorker = function(){
  if(children.worker === undefined){
    log('master', 'starting email worker');
    children.worker = cluster.fork({ROLE: 'worker'});
    children.worker.name = 'email worker';
    children.worker.on('online', function(){ log(children.worker, 'ready!'); });
    children.worker.on('exit', function(){
      log(children.worker, 'died :(');
      delete children.worker;
    });
    children.worker.on('message', function(message){
      log(children.worker, JSON.stringify(message));
    });
  }
};
```

...IT'S ALL JUST MESSAGE PASSING

## LOCAL MESSAGES

---

```
var doWorkerStuff = function(){
  process.on('message', function(message){
    emails.push(message);
  });

  var sendEmail = function(to, subject, text, callback){
    ...
  };

  var workerLoop = function(){
    if(emails.length === 0){
      setTimeout(workerLoop, 1000);
    }else{
      var e = emails.shift();
      process.send({msg: 'trying to send an email...'});
      sendEmail(e.to, e.subject, e.text, function(error){
        if(error){
          emails.push(e); // try again
          process.send({msg: 'failed sending email, trying again :('});
        }else{
          process.send({msg: 'email sent!'});
        }
      });
      setTimeout(workerLoop, 1000);
    }
  };

  workerLoop();
};
```

The diagram illustrates the flow of local messages and interprocess communication (IPC) in a Node.js worker process. It features four horizontal bars with arrows pointing to specific code locations:

- Message Queue:** Points to the `emails.push(message);` line in the `doWorkerStuff` function.
- Throttling:** Points to the `setTimeout(workerLoop, 1000);` line in the `workerLoop` function.
- Retry:** Points to the `emails.push(e); // try again` line in the `workerLoop` function.
- Interprocess Communication (IPC) with complex data-types:** Points to the `process.send({msg: 'email sent!'});` line in the `workerLoop` function.



**DEMO TIME**



# STRATEGY SUMMARY

### ▶ Notes:

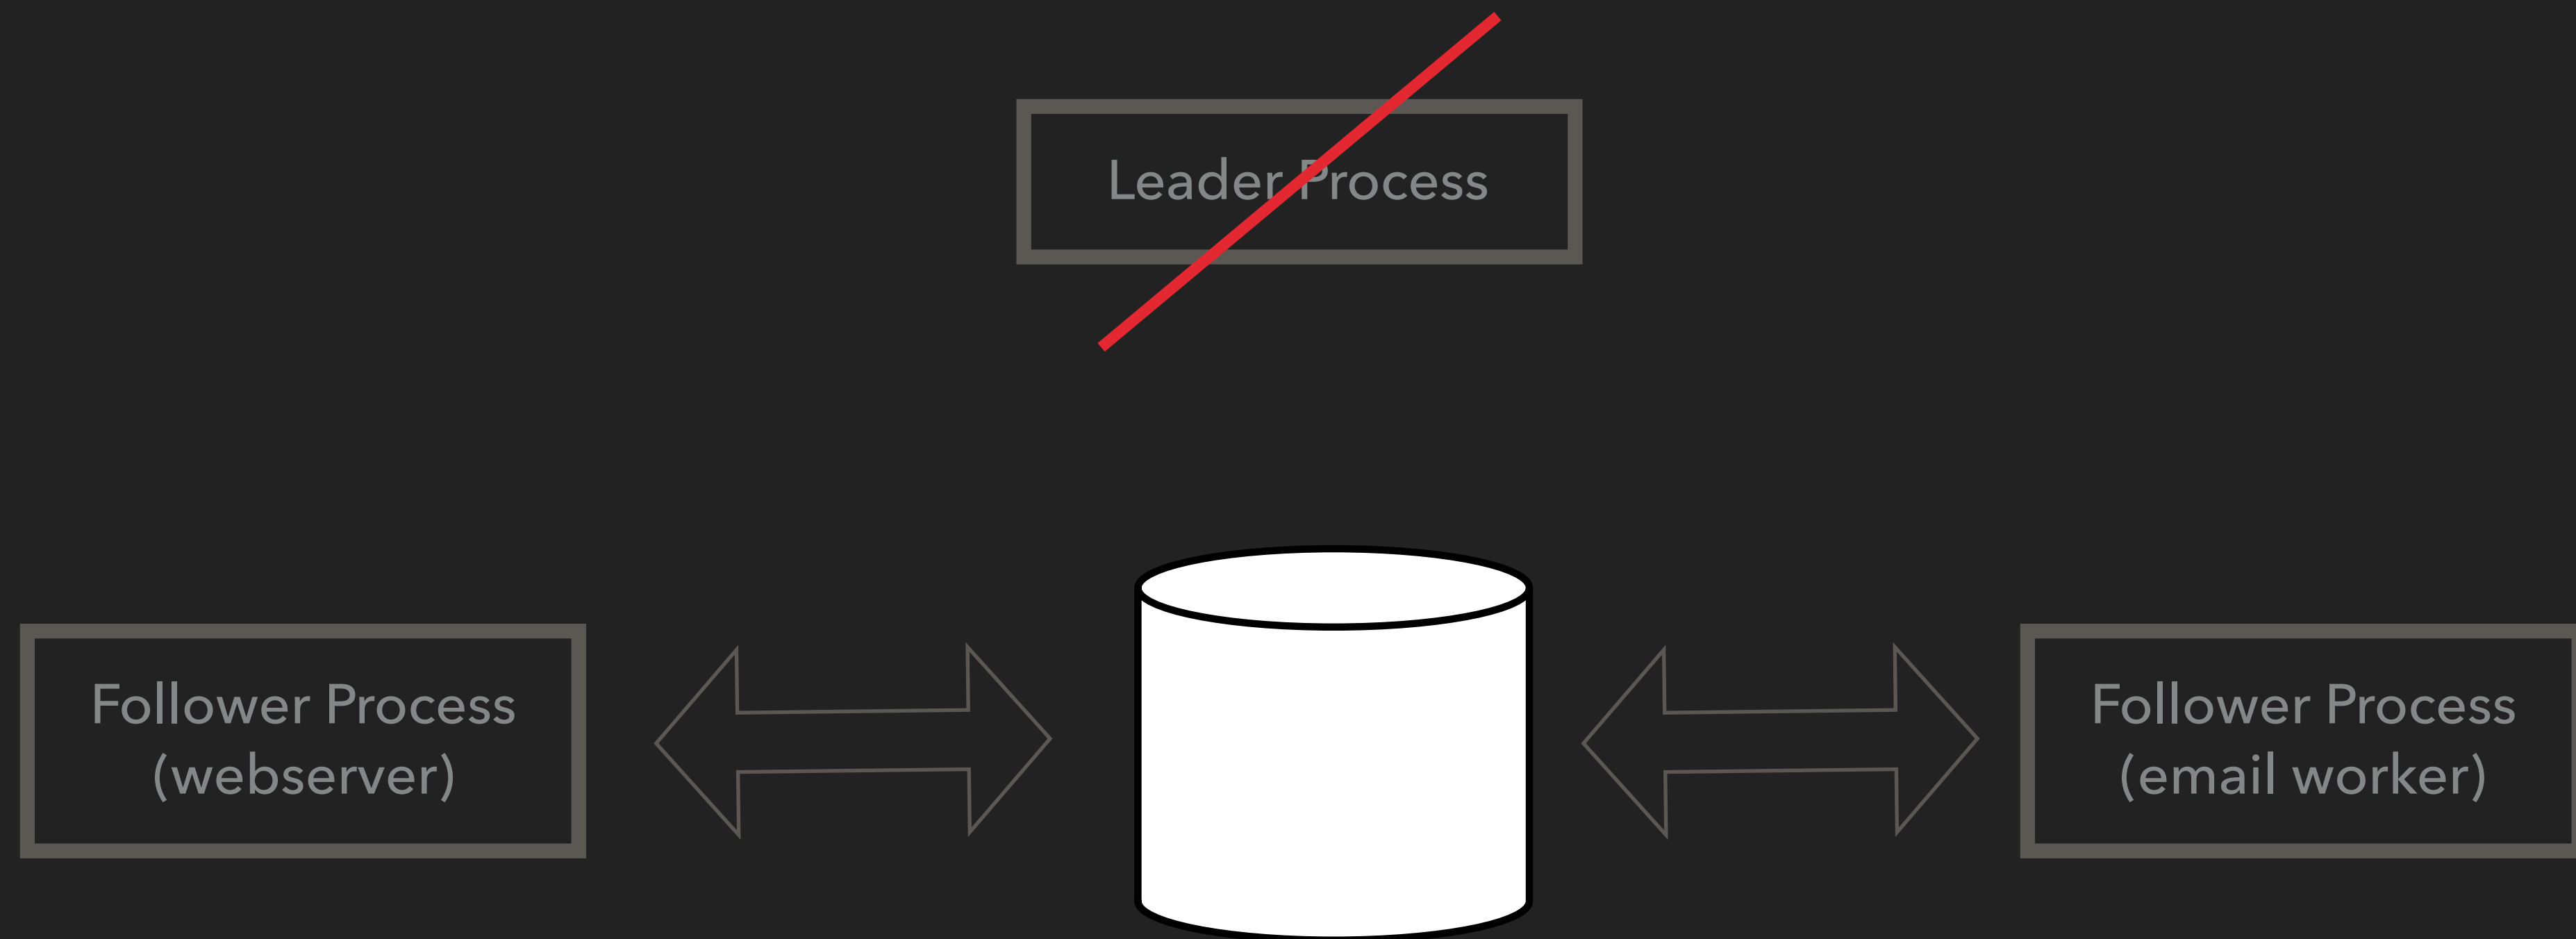
- ▶ the children never log themselves
  - ▶ the master does it for them
- ▶ Each process has it's own "main" loop:
  - ▶ web server
  - ▶ worker
  - ▶ master
- ▶ we can kill the child processes / allow them to crash...

## STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ In ~100 lines of JS...
    - ▶ Messages aren't lost when server dies
    - ▶ Web-server process not bothered by email sending
    - ▶ Error handling, Throttling, Queuing and retries!
    - ▶ Offline support?
- ▶ Why it is still a bad idea:
  - ▶ Bound to one server

# 4) REMOTE MESSAGES

## IMPROVEMENT IDEAS







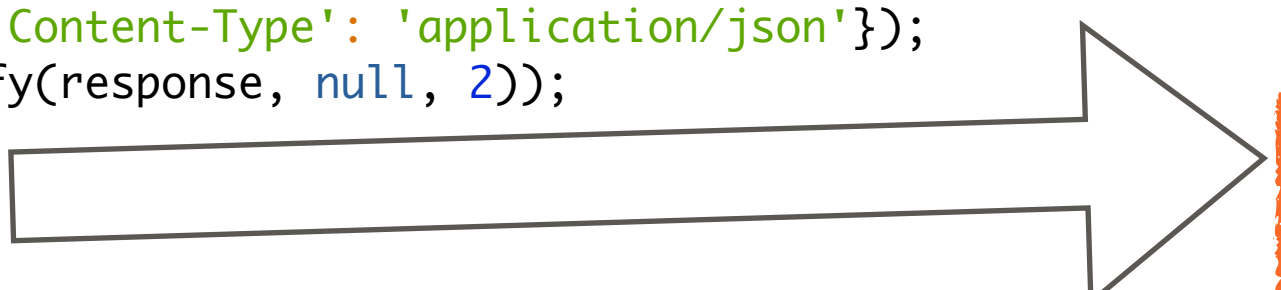
# IPC => REDIS PUB/SUB

```
var doServerStuff = function(){
  var server = function(req, res){
    var urlParts = req.url.split('/');
    var email = {
      to:      decodeURI(urlParts[1]),
      subject: decodeURI(urlParts[2]),
      text:    decodeURI(urlParts[3]),
    };

    var response = {email: email};
    res.writeHead(200, {'Content-Type': 'application/json'});
    res.end(JSON.stringify(response, null, 2));

    process.send(email);
  };

  http.createServer(server).listen(httpPort, '127.0.0.1');
};
```



```
var doServerStuff = function(){
  var publisher = Redis();
```

```
  var server = function(req, res){
    var urlParts = req.url.split('/');
    var email = {
      to:      decodeURI(urlParts[1]),
      subject: decodeURI(urlParts[2]),
      text:    decodeURI(urlParts[3]),
    };
  };
```

```
    publisher.publish(channel, JSON.stringify(email), function(){
      var response = {email: email};
      res.writeHead(200, {'Content-Type': 'application/json'});
      res.end(JSON.stringify(response, null, 2));
    });
  };
};
```

```
http.createServer(server).listen(httpPort, httpHost);
console.log('Server running at ' + httpHost + ':' + httpPort);
console.log('send an email and message to /TO_ADDRESS/SUBJECT/YOUR_MESSAGE');
};
```


## REMOTE MESSAGES

```
var doWorkerStuff = function(){
  process.on('message', function(message){
    emails.push(message);
  });

  var sendEmail = function(to, subject, text, callback){
    ...
  };

  var workerLoop = function(){
    if(emails.length === 0){
      setTimeout(workerLoop, 1000);
    }else{
      var e = emails.shift();
      process.send({msg: 'trying to send an email...'});
      sendEmail(e.to, e.subject, e.text, function(error){
        if(error){
          emails.push(e); // try again
          process.send({msg: 'failed sending email, trying again :('});
        }else{
          process.send({msg: 'email sent!'});
        }
      });
      setTimeout(workerLoop, 1000);
    }
  };
};

workerLoop();
};
```



```
var subscriber = Redis();

subscriber.subscribe(channel);
subscriber.on('message', function(channel, message){
  console.log('Message from Redis!');
  emails.push(JSON.parse(message));
});
```

Still with Throttling and Retry!



**DEMO TIME**



# STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ Redis Drivers are awesome
    - ▶ Message Buffering (for connection errors)
    - ▶ Thread-pools
    - ▶ Good language features (promises and callbacks)
  - ▶ Now we can use more than one server!
- ▶ Why it is still a bad idea:
  - ▶ Errors are logged, not passed back to the client
  - ▶ Email payload is lost on error or worker failure

# 5) REMOTE QUEUE

# IMPROVEMENT IDEAS

- ▶ Observability
  - ▶ How long is the queue?
  - ▶ How long does an item wait in the queue?
  - ▶ Operational Monitoring
- ▶ Redundancy
  - ▶ Backups
  - ▶ Clustering
  - ▶ Backups



### DATA STRUCTURES NEEDED FOR AN MVP QUEUE

- ▶ Array
  - ▶ push, pop, length

### DATA STRUCTURES NEEDED FOR A GOOD QUEUE


- ▶ Array
  - ▶ push, pop, length
- ▶ Hash (key types: string, integer, hash)
  - ▶ Set, Get, Exists
- ▶ Sorted Set
  - ▶ Exists, Add, Remove

REDIS HAS THEM ALL!



## RESQUE: DATA STRUCTURE FOR QUEUES IN REDIS

```
var queue = new NR.queue({connection: connectionDetails}, jobs);  
  
queue.on('error', function(error){ console.log(error); });  
  
queue.connect(function(){  
  queue.enqueue('math', "add", [1,2]);  
  queue.enqueueIn(3000, 'math', "subtract", [2,1]);  
});
```



# RESQUE: DATA STRUCTURE FOR QUEUES IN REDIS

redis Commander

Refresh

Commands

More...

127.0.0.1:6379:0

resque:\* (2)

queue:\* (1)

emailQueue (1)

queues (1)

Add New Value...

Delete Key

◀◀

◀

0

▶

▶▶

Goto Index

#	Value
0	<pre>{"class":"sendEmail","queue":"emailQueue","args":[{"to":"evantahler@gmail.com","subject":"hello_from_node","text":"hello_again"}]}</pre>

# RESQUE: DATA STRUCTURE FOR QUEUES IN REDIS

Overview

Working

Failed

Queues

Workers

Stats

Schedule

Delayed

## Queues

The list below contains all the registered queues with the number of jobs currently in the queue. Select a queue from above to view all jobs currently pending on the queue.

Name	Jobs
<u>emailQueue</u>	1
<u>failed</u>	0

## 0 of 0 Workers Working

The list below contains all workers which are currently running a job.

	Where	Queue	Processing
Nothing is happening right now...			



# USING NODE-RESQUE

```
var transporter = nodemailer.createTransport({
  service: 'gmail',
  auth: {
    user: require('./.emailUsername'),
    pass: require('./.emailPassword')
  }
});

var jobs = {
  sendEmail: function(data, callback){
    var email = {
      from:    require('./.emailUsername'),
      to:      data.to,
      subject: data.subject,
      text:    data.text,
    };

    transporter.sendMail(email, function(error, info){
      callback(error, {email: email, info: info});
    });
  }
};
```

SENDING EMAILS IS A “JOB” NOW

# USING NODE-RESQUE

```
var server = function(req, res){
  var urlParts = req.url.split('/');
  var email    = {
    to:        decodeURI(urlParts[1]),
    subject:   decodeURI(urlParts[2]),
    text:      decodeURI(urlParts[3]),
  };
};
```

```
queue.enqueue('emailQueue', "sendEmail", email, function(error){
  if(error){ console.log(error) }
  var response = {email: email};
  res.writeHead(200, {'Content-Type': 'application/json'});
  res.end(JSON.stringify(response, null, 2));
});
};
```

```
var queue = new NR.queue({connection: connectionDetails, jobs});
queue.connect(function(){
  http.createServer(server).listen(httpPort, httpHost);
  console.log('Server running at ' + httpHost + ':' + httpPort);
  console.log('send an email and message to /TO_ADDRESS/SUBJECT/YOUR_MESSAGE');
});
```

# USING NODE-RESQUE

```
var worker = new NR.worker({connection: connectionDetails, queues: ['emailQueue']}, jobs);
worker.connect(function(){
  worker.workerCleanup();
  worker.start();
});
```

```
worker.on('start', function(){ console.log("worker started"); });
worker.on('end', function(){ console.log("worker ended"); });
worker.on('cleaning_worker', function(worker, pid){ console.log("cleaning old worker " + worker); });
worker.on('poll', function(queue){ console.log("worker polling " + queue); });
worker.on('job', function(queue, job){ console.log("working job " + queue + " " + JSON.stringify(job)); });
worker.on('reEnqueue', function(queue, job, plugin){ console.log("reEnqueue job (" + plugin + ") " + queue + " " + JSON.stringify(job)); });
worker.on('success', function(queue, job, result){ console.log("job success " + queue + " " + JSON.stringify(job) + " >> " + result); });
worker.on('failure', function(queue, job, failure){ console.log("job failure " + queue + " " + JSON.stringify(job) + " >> " + failure); });
worker.on('error', function(queue, job, error){ console.log("error " + queue + " " + JSON.stringify(job) + " >> " + error); });
worker.on('pause', function(){ console.log("worker paused"); });
```



**DEMO TIME**





**BUT WHAT IS SO SPECIAL ABOUT NODE.JS HERE?**

# IMPROVEMENT IDEAS

- ▶ The node.js event loops is great for processing all non-blocking events, not just web servers.
- ▶ Most Background jobs are non-blocking events
  - ▶ Update the DB, Talk to this external service, etc
- ▶ So node can handle many of these at once per process!
- ▶ Redis is fast enough to handle many "requests" from the same process in this manner
  - ▶ We can use the same connection or thread-pool



# USING NODE-RESQUE AND MAXIMIZING THE EVENT LOOP

```
var multiWorker = new NR.multiWorker({  
  connection: connectionDetails,  
  queues: ['slowQueue'],  
  minTaskProcessors: 1,  
  maxTaskProcessors: 20,  
}, jobs);
```

non-blocking

```
var jobs = {  
  "slowSleepJob": {  
    plugins: [],  
    pluginOptions: {},  
    perform: function(callback){  
      var start = new Date().getTime();  
      setTimeout(function(){  
        callback(null, (new Date().getTime() - start) );  
      }, 1000);  
    },  
  },  
};
```

blocking

```
  "slowCPUJob": {  
    plugins: [],  
    pluginOptions: {},  
    perform: function(callback){  
      var start = new Date().getTime();  
      blockingSleep(1000);  
      callback(null, (new Date().getTime() - start) );  
    },  
  },  
};
```

# HOW CAN YOU TELL IF THE EVENT LOOP IS BLOCKED?

// inspired by <https://github.com/tj/node-blocked>

```
module.exports = function(limit, interval, fn) {  
  var start = process.hrtime();  
  
  setInterval(function(){  
    var delta = process.hrtime(start);  
    var nanosec = delta[0] * 1e9 + delta[1];  
    var ms = nanosec / 1e6;  
    var n = ms - interval;  
    if (n > limit){  
      fn(true, Math.round(n));  
    }else{  
      fn(false, Math.round(n));  
    }  
    start = process.hrtime();  
  }, interval).unref();  
};
```

... SEE HOW LONG IT TAKES FOR THE NEXT “LOOP”



**DEMO TIME**



# REDIS

- ▶ Redis has unique properties that make it perfect for this type of workload
  - ▶ FAST
  - ▶ Single-threaded so you can have real array operations (pop specifically)
  - ▶ Data-structure creation on the fly (new queues)
  - ▶ Dependent on only RAM and network

## STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ In addition to persistent storage and multiple server/process support, you get get CPU scaling and Throttling very simply!
  - ▶ Integrates well with the resque/sidekiq ecosystem
- ▶ This is now a good idea!

# 6) IMMUTABLE EVENT BUS


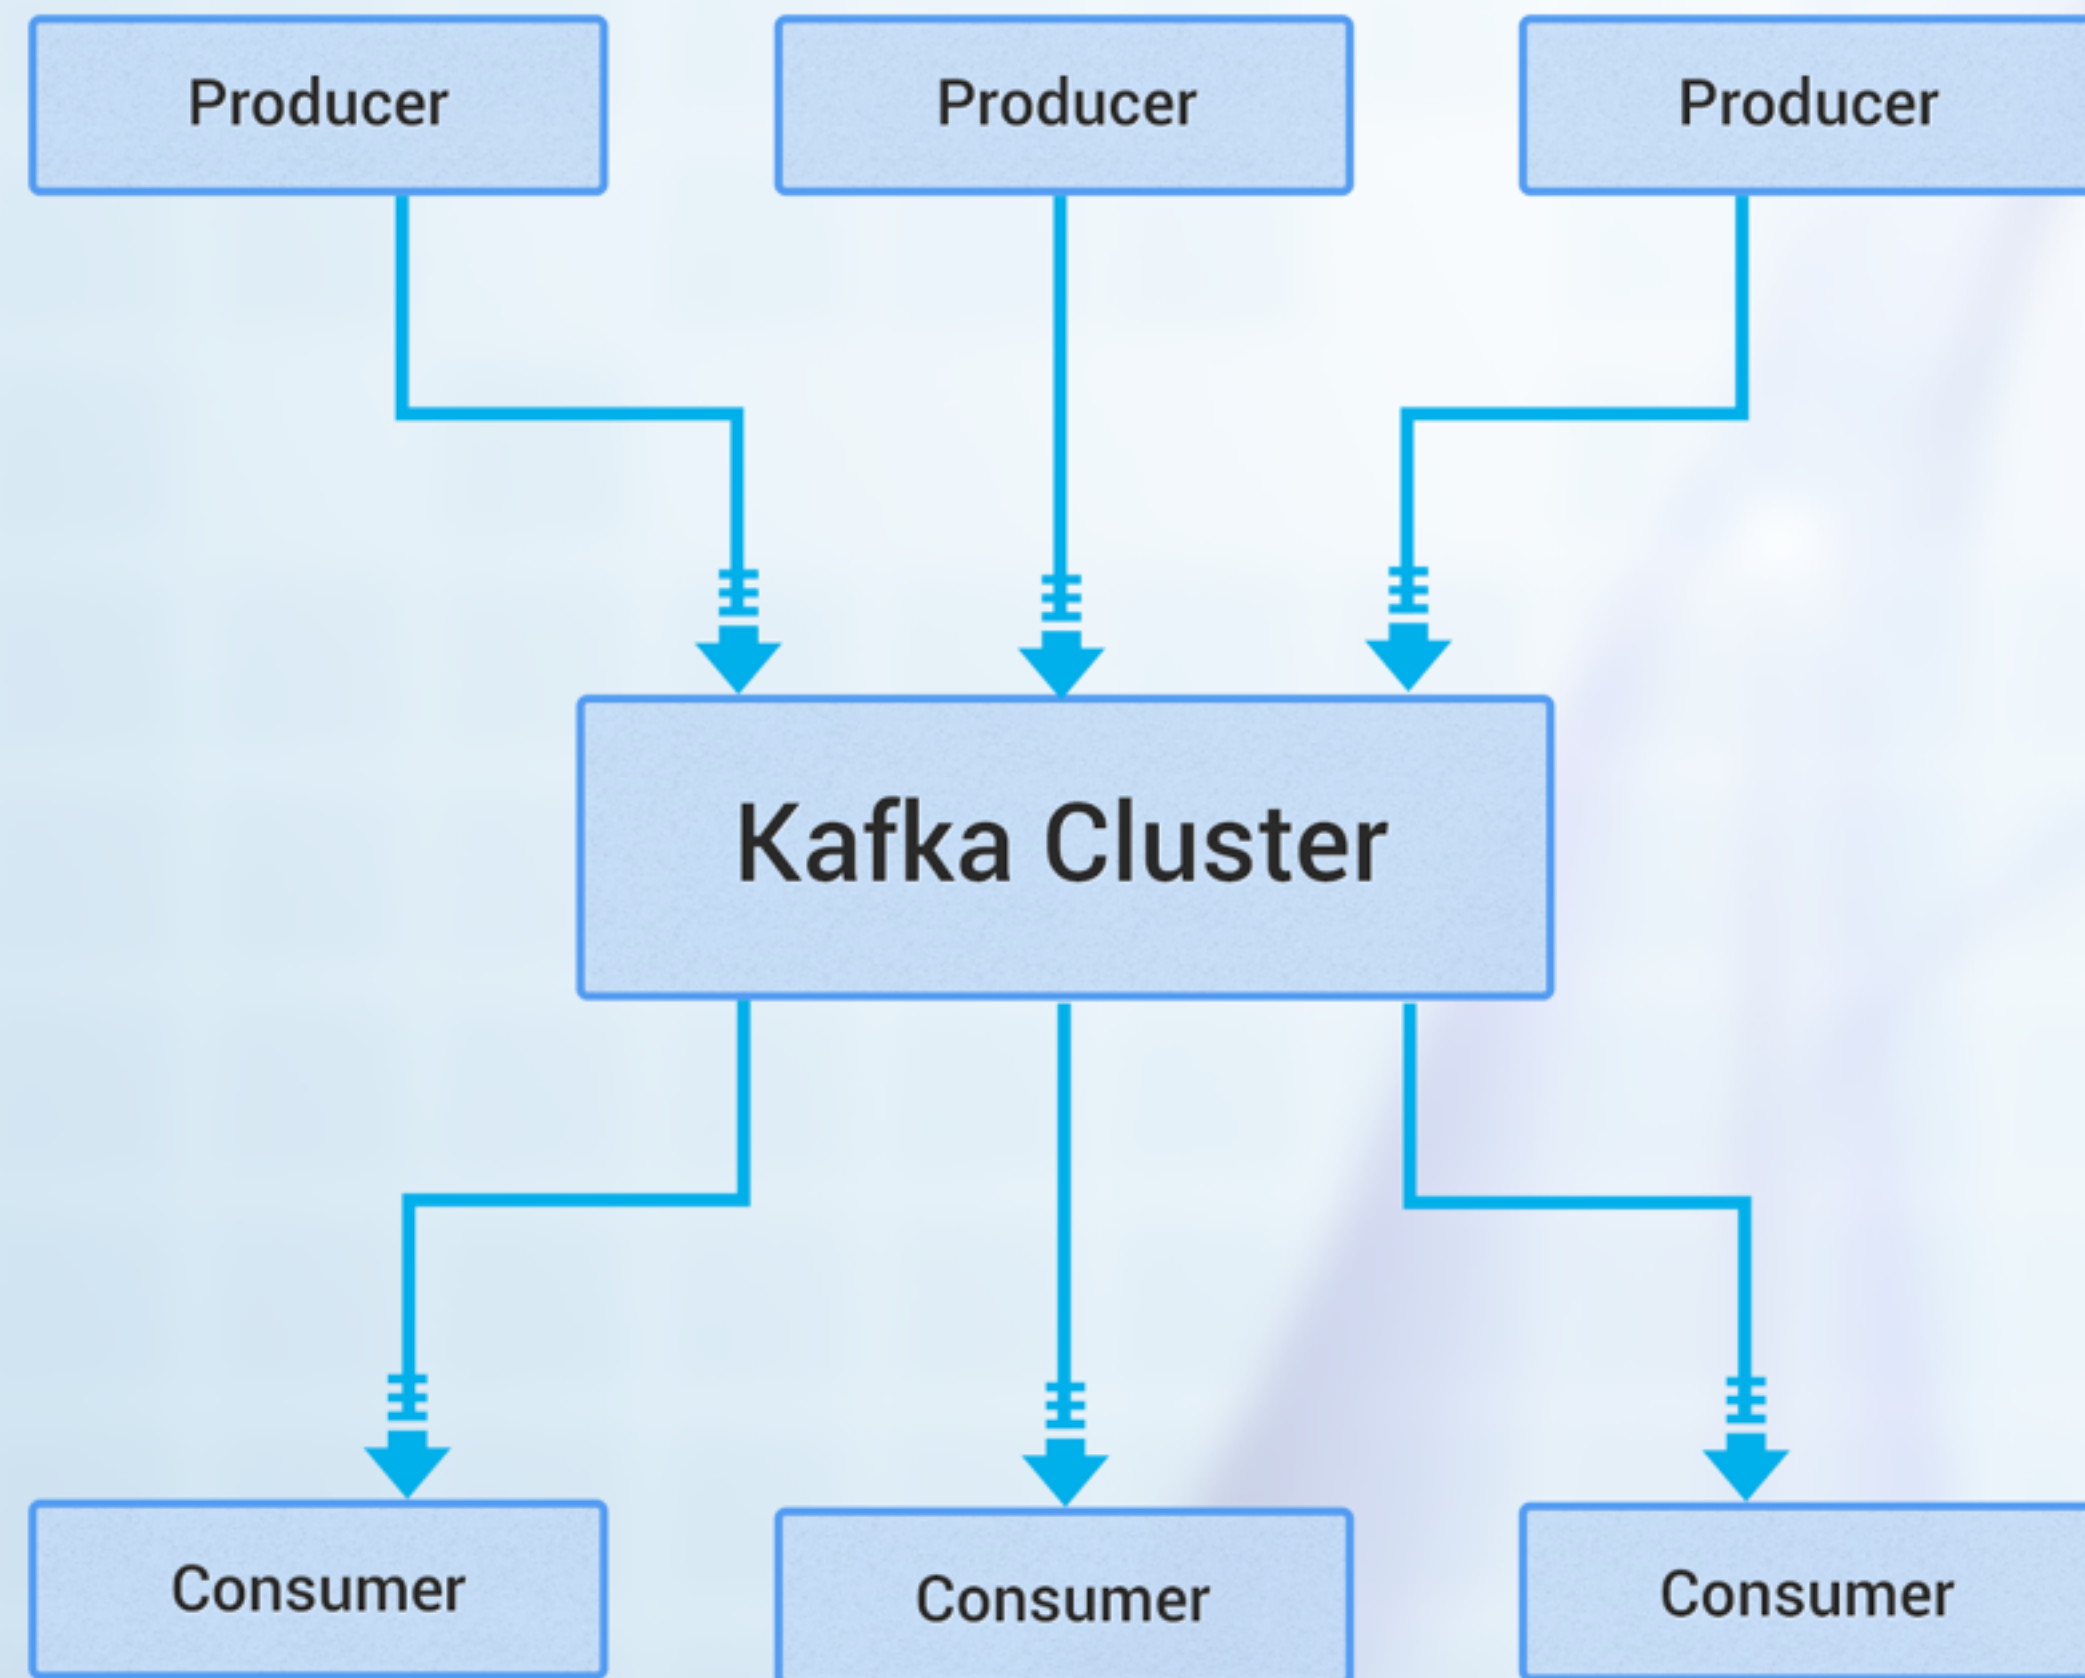
*"The Future"... "Maybe"*



# IMPROVEMENT IDEAS

- ▶ What was wrong with the resque pattern?
  - ▶ Jobs are consumed and deleted... no historical introspection
  - ▶ In redis, storing more and more events to a single key gains nothing from redis-cluster
  - ▶ If you wanted to do more than one thing on user#create, you would need to fire many jobs
- ▶ What if we just fired a "user\_created" event and let the workers choose what to do?

# IMMUTABLE EVENT BUS



- ▶ Events are written to a list and never removed
- ▶ Consumers know where their last read was and can continue

# IMPROVEMENT IDEAS

- ▶ What to we need that redis cannot do natively
  - ▶ A “blocking” get and incr
  - ▶ Tools to seek for “the next key” for listing partitions

... GOOD THING WE HAVE LUA!



# WHAT WOULD LUA LOOK LIKE FOR THIS?

```
var luaLines = [];  
// get the counter for this named consumer  
luaLines.push('local counter = 0');  
luaLines.push('if redis.call("HEXISTS", "'" + this.prefix + 'counters"', KEYS[1]) == 1 then');  
luaLines.push('  local counter = redis.call("HGET", "'" + this.prefix + 'counters"', KEYS[1])');  
luaLines.push('end');  
// if the partition exists, get the data from that key  
// otherwise + the partition and end  
luaLines.push('local partition = "'" + this.prefix + 'partitions:" .. KEYS[2]');  
luaLines.push('if reids.call("EXISTS", partition) == 0 then');  
luaLines.push('  return nil');  
luaLines.push('else ');  
luaLines.push('  local event = reids.call("LRANGE", partition, counter, counter)');  
luaLines.push('  reids.call("HSET", "'" + this.prefix + ':counters"', (counter + 1))');  
luaLines.push('  return event');  
luaLines.push('end');  
  
this.redis.defineCommand('getAndIncr', {  
  numberOfKeys: 2, lua: luaLines.join('\r\n')  
});
```

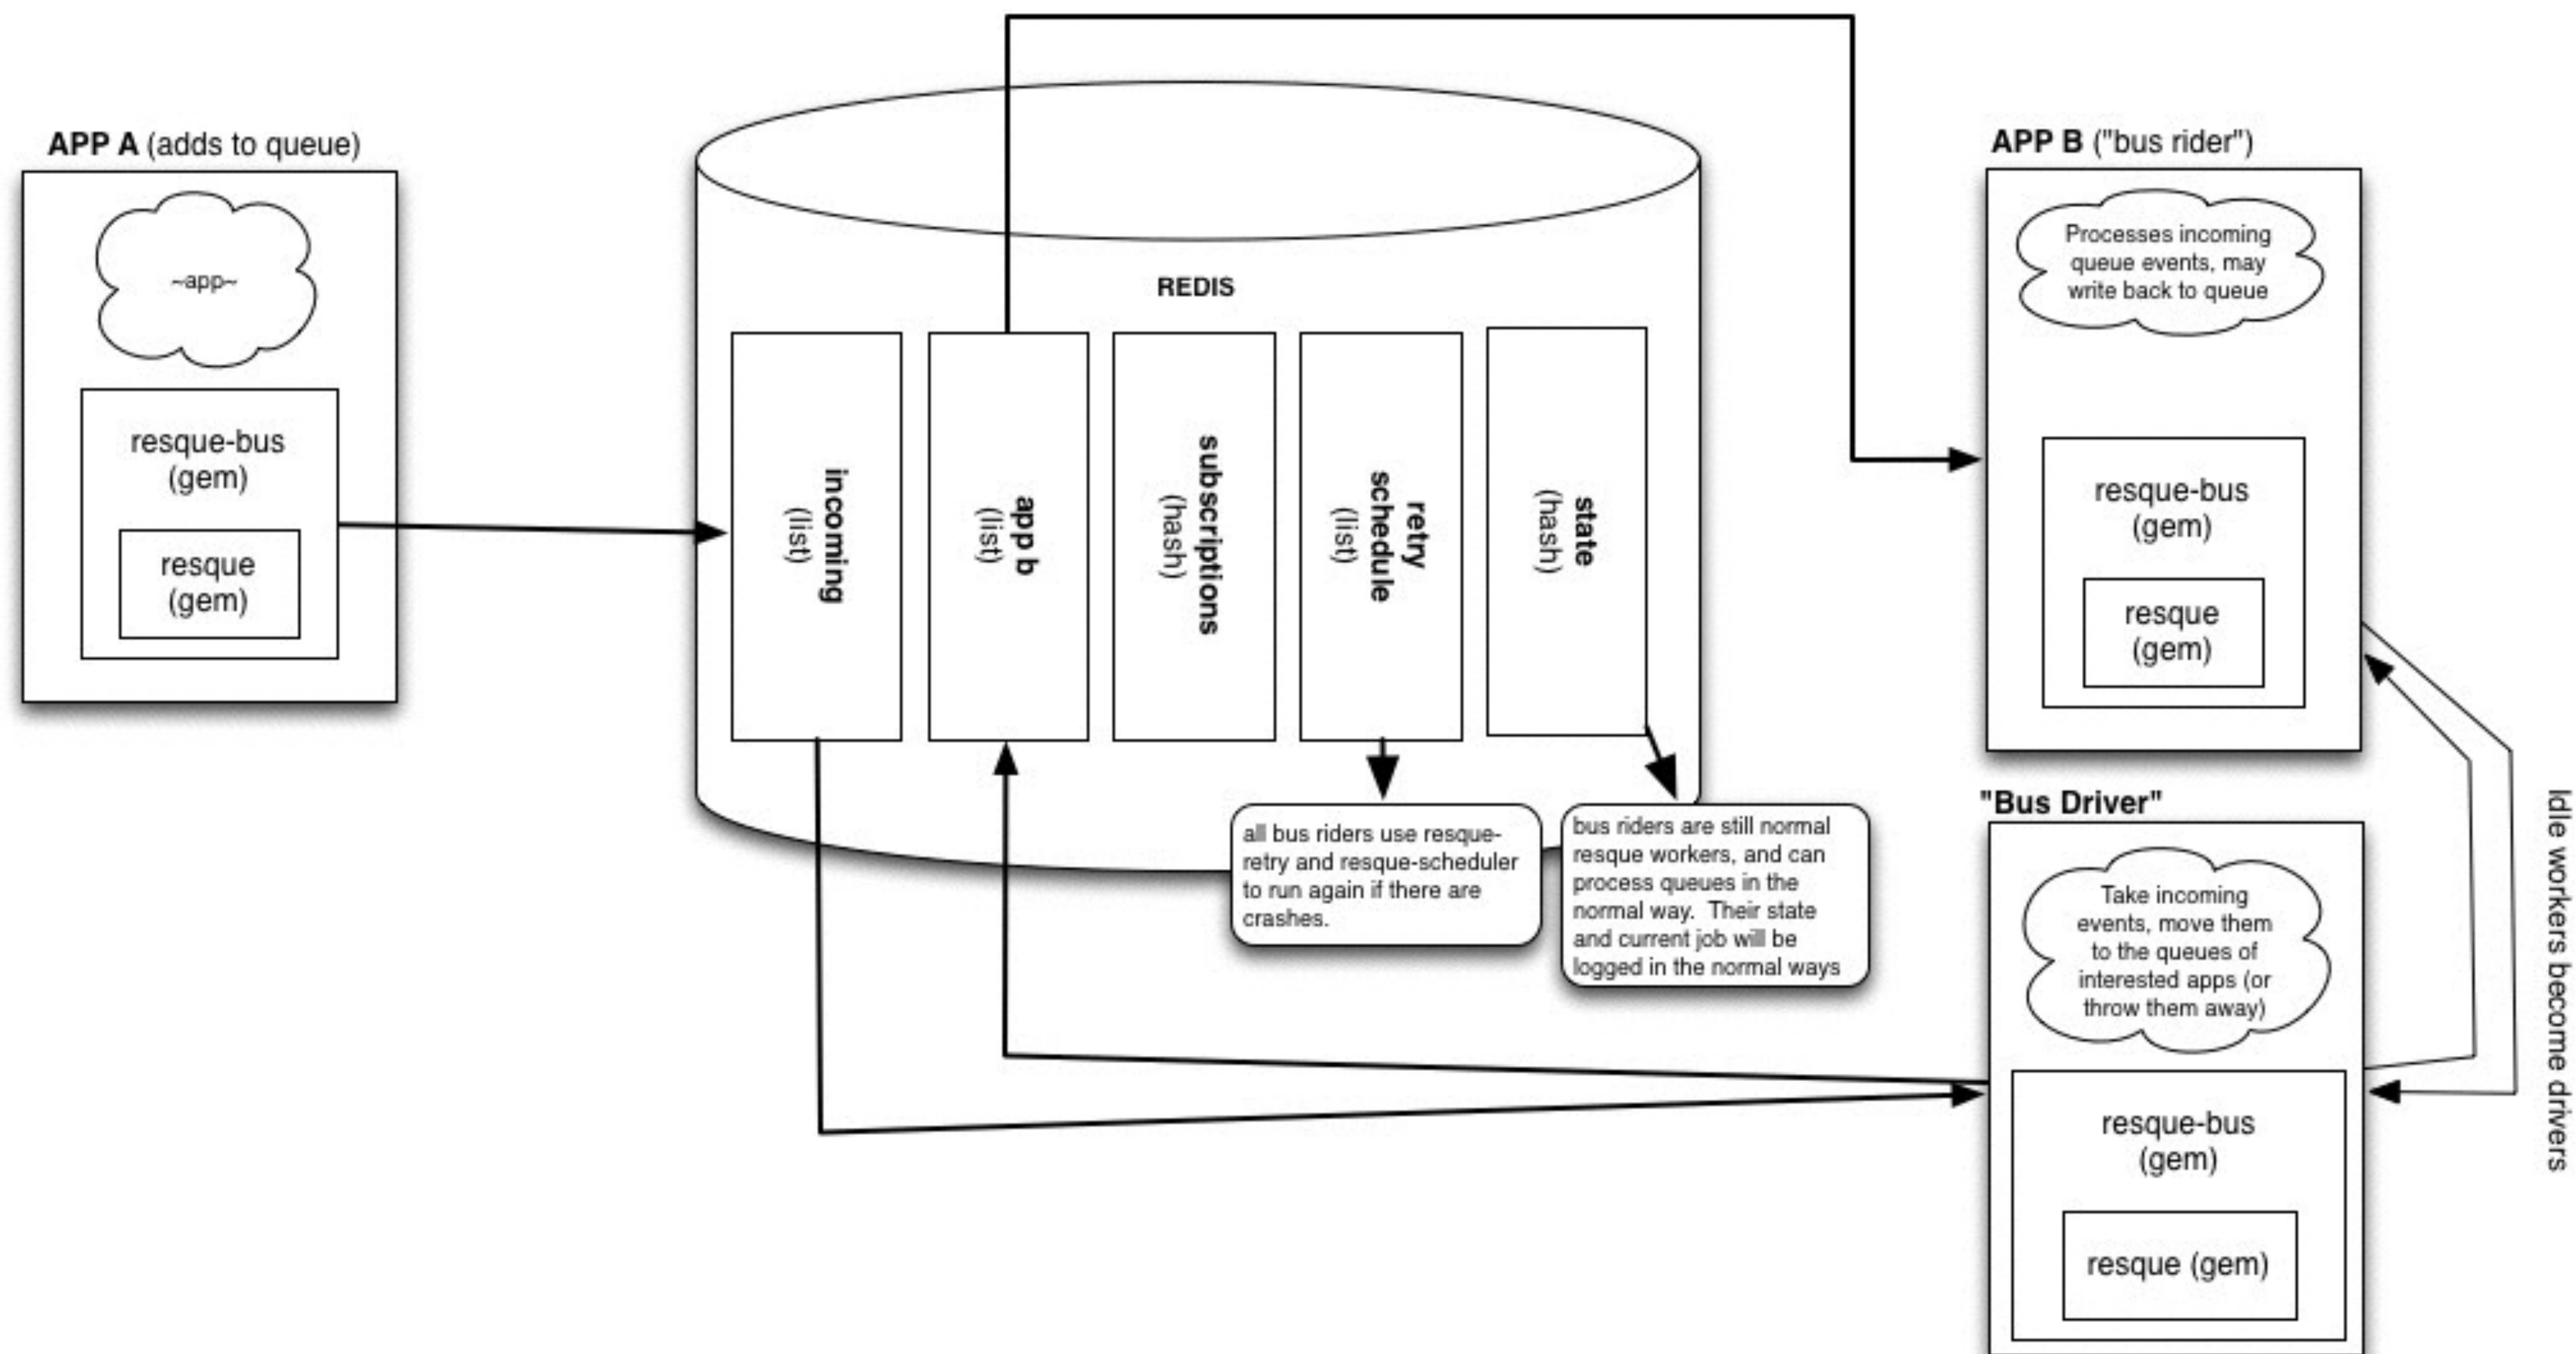
## STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ Just like with Resque, we can poll/stream many events at once (non-blocking)
- ▶ Why us this a bad idea?
  - ▶ We would need a *\*lot\** of LUA. We are actively slowing down Redis to preform more operations in a single block. With Cluster, we might have enough RAM for this, but this is more expensive than using another technology which stores data on disk.



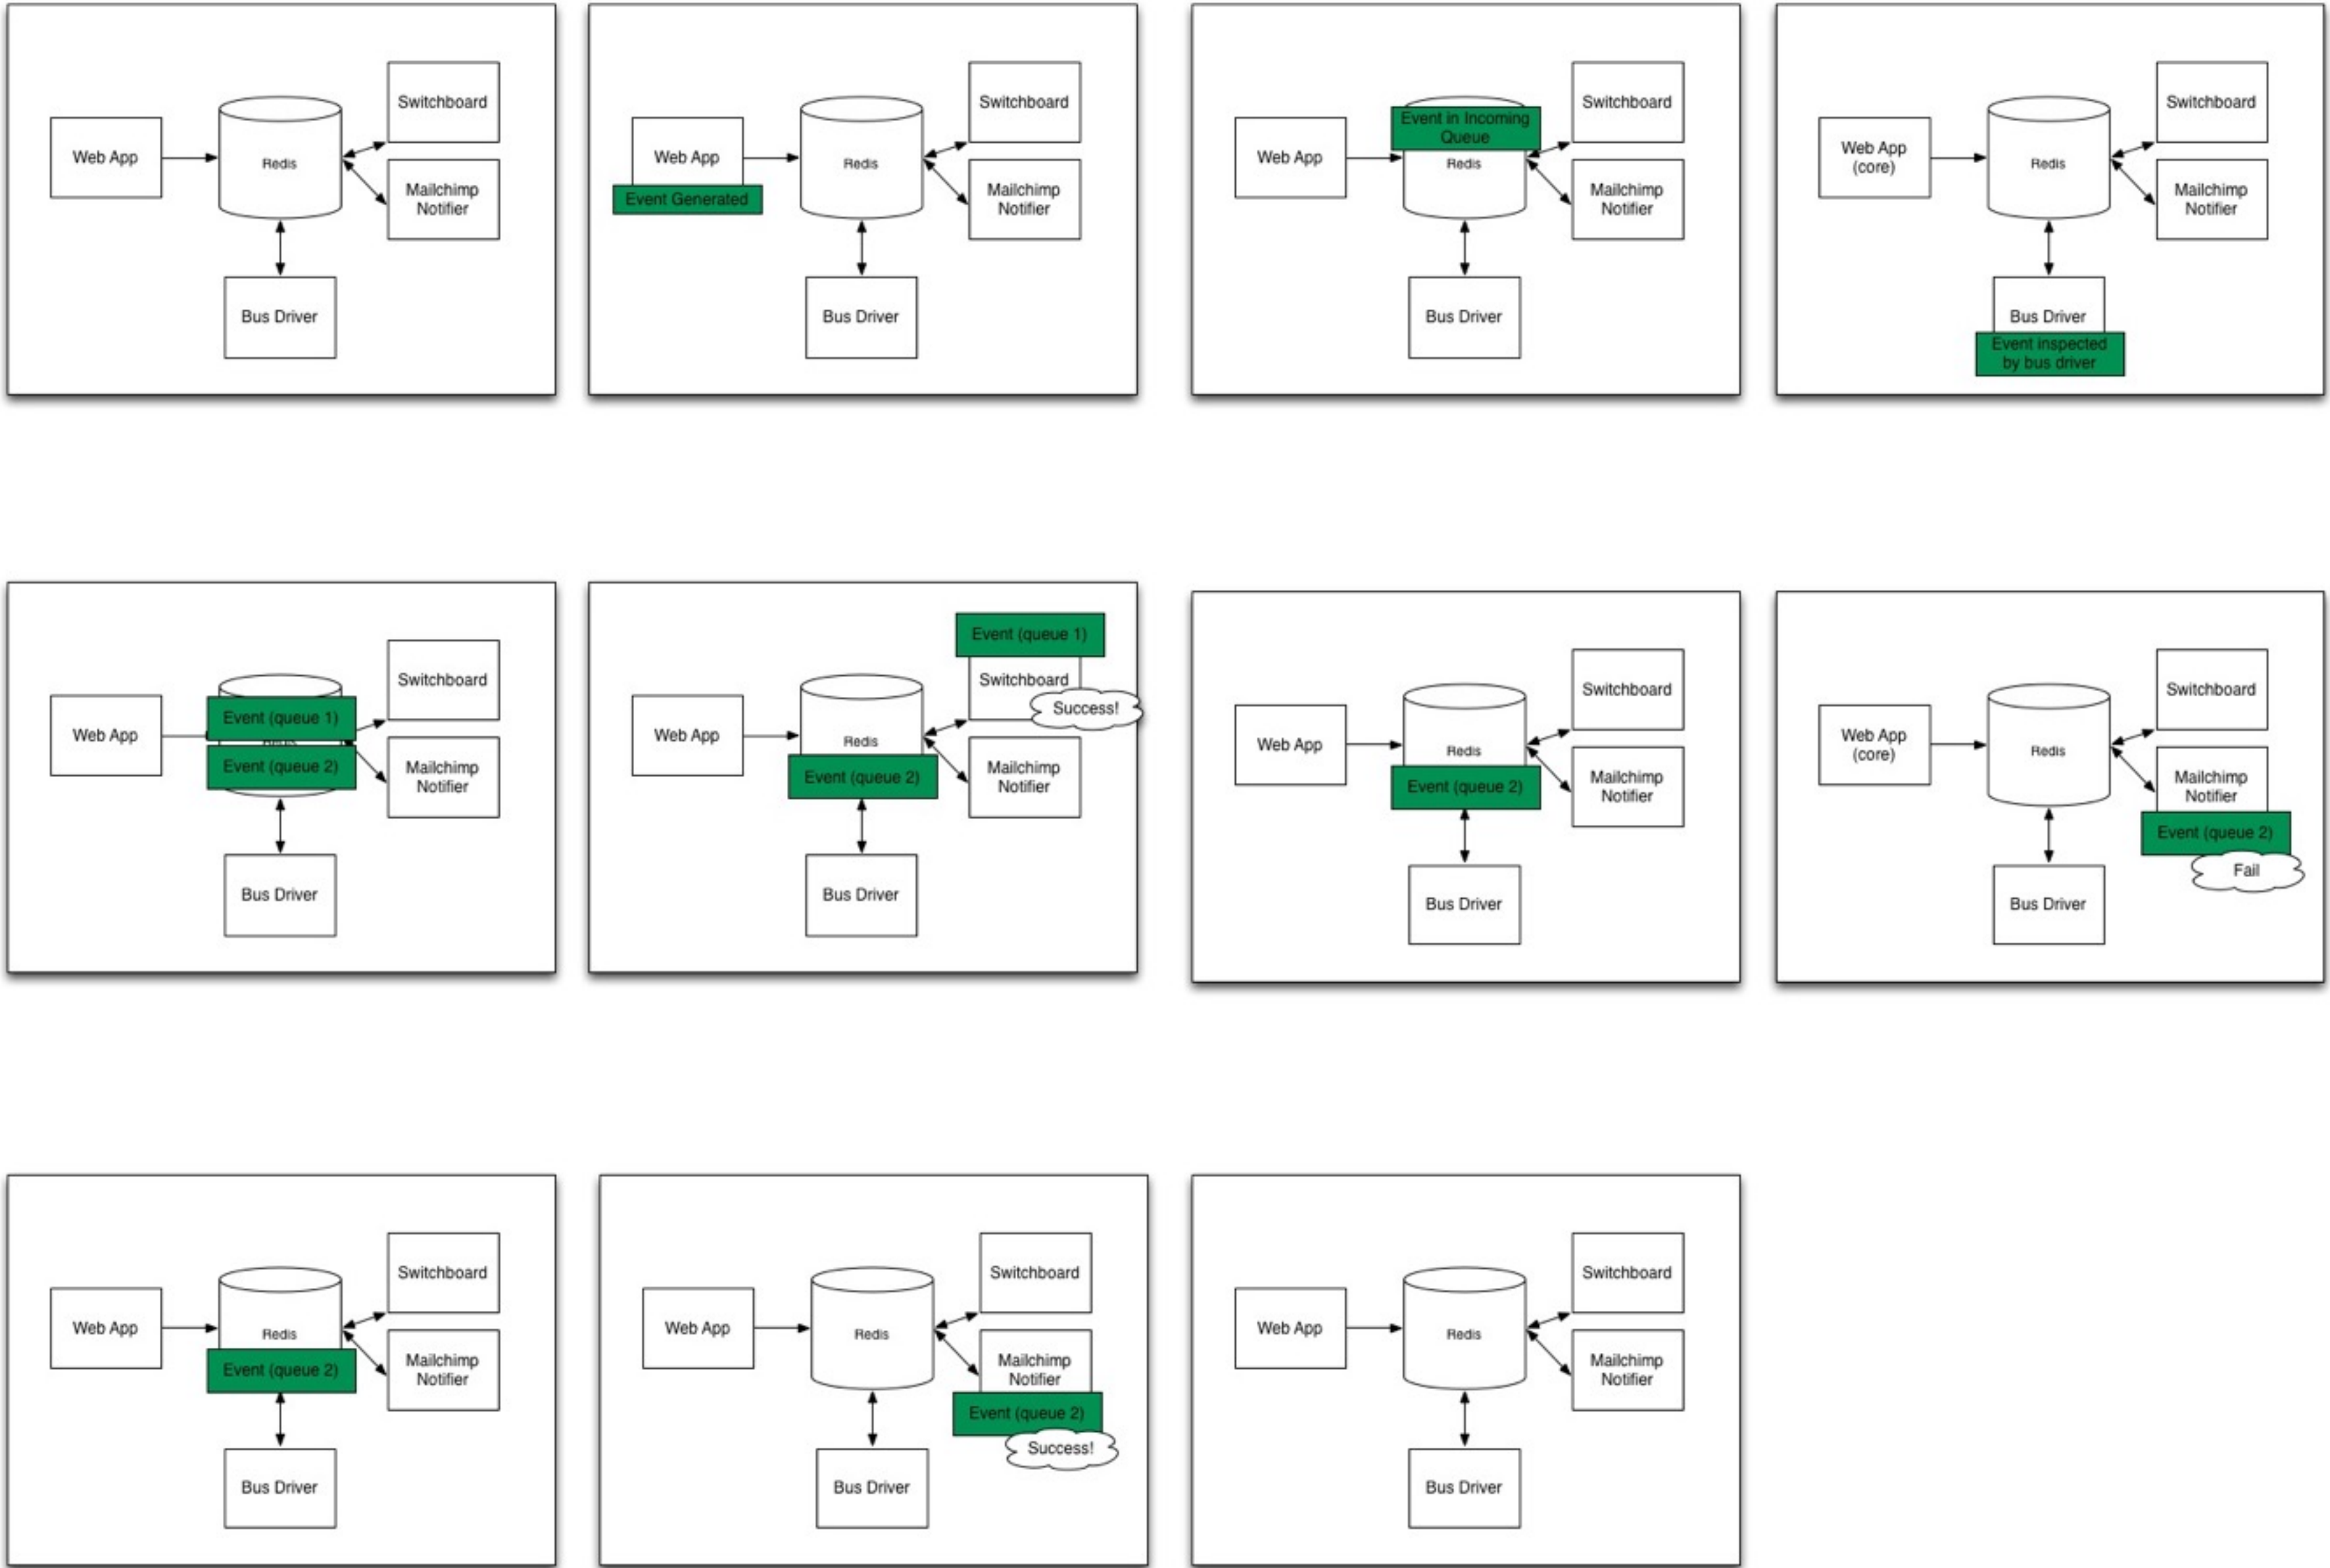
**CAN WE WE MEET IN THE MIDDLE  
OF “IMMUTABLE EVENT BUS” +  
“RESQUE”?**

# QUEUEBUS





# QUEUEBUS



## STRATEGY SUMMARY

- ▶ Why it is better in node:
  - ▶ Just like with Resque, we can poll/stream many events at once (non-blocking)
- ▶ What did we gain on Resque?
  - ▶ Syndication of events to multiple consumers
- ▶ What didn't we get from Kafka?
  - ▶ repayable event log



# BACKGROUND TASKS IN NODE.JS



**NODE-RESQUE:**

<https://github.com/taskrabbit/node-resque>

**SUPPORTING PROJECT:**

[https://github.com/evantahler/background\\_jobs\\_node](https://github.com/evantahler/background_jobs_node)

**QUEUE-BUS:**

<https://github.com/queue-bus>

**THESE SLIDES:**

<http://bit.ly/1pDm9Vf>

**@EVANTAHLER**