
Background Jobs + Node.JS

Async Processing for your Async Language

{{ BioPage }}

@evantahler

www.evantahler.com



- Director of Technology @ [TaskRabbit](#)
- Maintainer of [node-resque](#)
- Maintainer of [actionhero.js](#)

We are Hiring! Talk to me later!

DISCLAIMER!

Most of what you will see is a terrible idea.

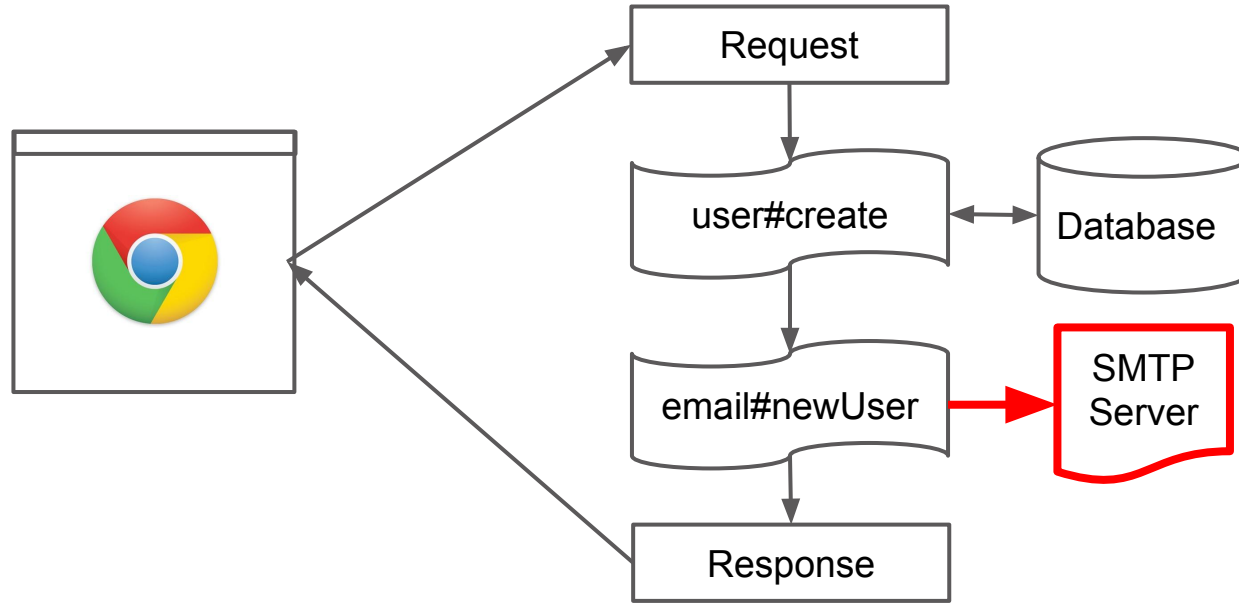
Try this at ~, not on production

The Point:

Everything is
better/faster/stronger in node.

...even the bad ideas!

So you have a website...



Possible Task Strategies

1. Foreground (in-line)
 2. Parallel (threaded-ish)
 3. Local Messages (fork-ish)
 4. ~~Remote Messages~~
 5. Remote Queues (Resque-ish)
 6. ~~Event Bus (Kafka-ish)~~
-

Strategy 1: Foreground

php

Downloads

Documentation

Get Involved

Help

```
<?php
$to      = 'nobody@example.com';
$subject = 'the subject';
$message = 'hello';
$headers = 'From: webmaster@example.com' . "\r\n" .
    'Reply-To: webmaster@example.com' . "\r\n" .
    'X-Mailer: PHP/' . phpversion();

mail($to, $subject, $message, $headers);
?>
```

Strategy 1: Foreground

```
var http      = require('http');  
var nodemailer = require('nodemailer');  
var httpPort  = 8080 || process.env.port;  
  
var transporter = nodemailer.createTransport({  
  service: 'gmail',  
  auth: { user: require('./.emailUsername'), pass: require('./.emailPassword') }  
});
```



```
var server = function(req, res){
  var start = Date.now();
  var responseCode = 200;
  var response = {};
  sendEmail(req, function(error, email){
    response.email = email;
    if(error){
      console.log(error);
      responseCode = 500;
      response.error = error;
    }
    res.writeHead(responseCode, {'Content-Type': 'application/json'});
    res.end(JSON.stringify(response, null, 2));
    var delta = Date.now() - start;
    console.log('Sent an email to in ' + delta + 'ms');
  });
};
```

```
http.createServer(server).listen(httpPort, '0.0.0.0');
```

```
console.log('Server running at http://0.0.0.0:' + httpPort);
console.log('send an email and message to /TO_ADDRESS/SUBJECT/YOUR_MESSAGE');
```

Strategy 1: Foreground

```
var sendEmail = function(req, callback){  
  var urlParts    = req.url.split('/');  
  var email       = {  
    from:    require('./emailUsername'),  
    to:      urlParts[1],  
    subject: urlParts[2],  
    text:    urlParts[3],  
  };  
  transporter.sendMail(email, function(error, info){  
    callback(error, email);  
  });  
};
```

Async!



Strategy 1: Foreground

- Why it is better in node:
 - The client still needs to wait for the message to send, but you won't block any other client's requests
 - Avg response time of ~2 seconds from my couch
 - Why it is still a bad idea:
 - Slow
 - Spending “web server” resources on sending email
 - Error / Timeout to the client for “partial success”
 - IE: Account created but email not sent
 - Confusing to the user, dangerous for the DB
-

Strategy 2: Parallel

- “Threading”
 - But if it were real threading, the client would still have to wait
 - I guess this might help you catch errors...
 - But you could use domains?
 - **note: do not get into a discussion about threads*
 - Lets get crazy:
 - **Ignore the Callback**
-

DISCLAIMER!

Most of what you will see is a terrible idea.

Try this at ~, not on production

```
var sendEmail = function(req, callback){  
  var urlParts    = req.url.split('/');  
  var email       = {  
    from:    require('./emailUsername'),  
    to:      urlParts[1],  
    subject: urlParts[2],  
    text:    urlParts[3],  
  };  
  transporter.sendMail(email, function(error, info){  
    if(typeof callback === 'function'){ callback(error, email); }  
  });  
};
```

Async!



```
var server = function(req, res){  
  var start = Date.now();  
  var responseCode = 200;  
  var response     = {};  
  sendEmail(req);  
  res.writeHead(responseCode, {'Content-Type': 'application/json'});  
  res.end(JSON.stringify(response, null, 2));  
};
```

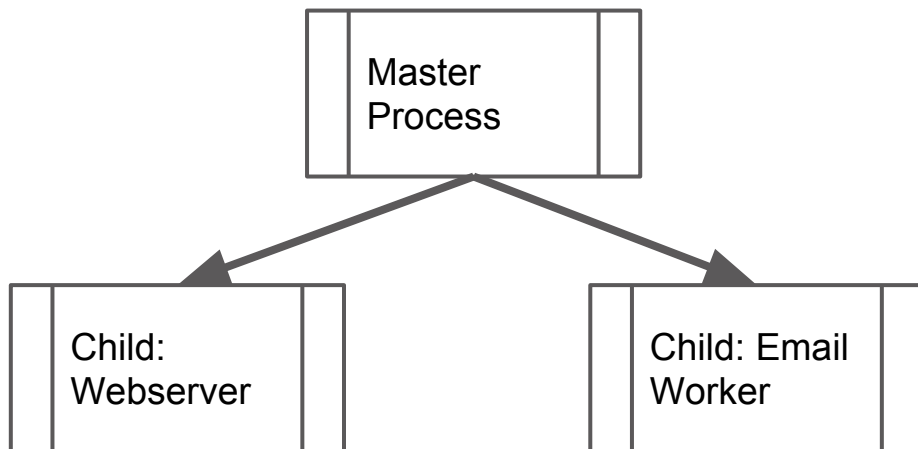


Strategy 2: Parallel

- Why it is better in node:
 - It's rare you can actually do this in a language... without threading!
 - Crazy-wicked-fast.
 - Why it is still a bad idea:
 - 0 callbacks, 0 data captured
 - I guess you could log errors?
 - But what would you do with that data?
 - The client has no idea what happened
-


Strategy 3: Local Messages

- “Forking”
 - or: *“The part of the talk where we grossly over-engineer some stuff”*



Strategy 3: Local Messages

```
if(cluster.isMaster){  
  doMasterStuff();  
}else{  
  if(process.env.ROLE === 'server'){ doServerStuff(); }  
  if(process.env.ROLE === 'worker'){ doWorkerStuff(); }  
}
```

```
var doServerStuff = function(){  
  var server = function(req, res){  
    var urlParts = req.url.split('/');  
    var email = {  
      to: urlParts[1],  
      subject: urlParts[2],  
      text: urlParts[3]  
    };  
  
    var response = {email: email};  
    res.writeHead(200, {'Content-Type': 'application/json'});  
    res.end(JSON.stringify(response, null, 2));  
  
    process.send(email);  IPC!  
  };  
  
  http.createServer(server).listen(httpPort);  
};
```

```
var doMasterStuff = function(){  
    log('master', 'started master');  
  
    var masterLoop = function(){  
        checkOnWebServer();  
        checkOnEmailWorker();  
    };  
  
    var checkOnWebServer = function() { 🚚 };  
  
    var checkOnEmailWorker = function() { 🚚 };  
    |  
    setInterval(masterLoop, 1000);  
};
```

```
var checkOnWebServer = function(){
  if(children.server === undefined){
    log('master', 'starting web server');
    children.server = cluster.fork({ROLE: 'SERVER'});
    children.server.name = 'web server';
    children.server.on('online',    function(){ log(children.server, 'ready on port ' + httpPort); });
    children.server.on('exit',      function(){
      log(children.server, 'died :(');
      delete children.server;
    });
    children.server.on('message',   function(message){
      log(children.server, 'got an email to send from the webserver: ' + JSON.stringify(message));
      children.worker.send(message);
    });
  }
};
```

**It's really all just
message passing and
monitoring...**

```
var checkOnEmailWorker = function(){
  if(children.worker === undefined){
    log('master', 'starting email worker');
    children.worker = cluster.fork({ROLE: 'WORKER'});
    children.worker.name = 'email worker';
    children.worker.on('online',    function(){ log(children.worker, 'ready!'); });
    children.worker.on('exit',      function(){
      log(children.worker, 'died :(');
      delete children.worker;
    });
    children.worker.on('message',  function(message){
      log(children.worker, JSON.stringify(message));
    });
  }
};
```

**It's really all just
message passing and
monitoring...**

```
var doWorkerStuff = function(){  
  process.on('message', function(message){  
    emails.push(message);  
  });  
};
```

Message Queue!

```
var sendEmail = function(to, subject, text, callback){  
  var email = {  
    from: require('./emailUsername'),  
    to: to,  
    subject: subject,  
    text: text,  
  };  
};
```

```
  transporter.sendMail(email, function(error, info){  
    callback(error, email);  
  });  
};
```

```
var workerLoop = function(){  
  if(emails.length === 0){  
    setTimeout(workerLoop, 1000);  
  }else{  
    var e = emails.shift();  
    process.send({msg: 'trying to send an email...'});  
    sendEmail(e.to, e.subject, e.text, function(error){  
      if(error){  
        emails.push(e); // try again  
        process.send({msg: 'failed sending email, trying again :('});  
      }else{  
        process.send({msg: 'email sent!'});  
      }  
      setTimeout(workerLoop, 1000);  
    });  
  }  
};  
  
workerLoop();  
};
```

Throttling!

Retry!

IPC!

Strategy 3: Local Messages

- Notes:
 - the children never log themselves
 - the master does it for them
 - Each process has it's own “main” loop:
 - web server
 - worker
 - master
 - AND we can kill the child processes...
-

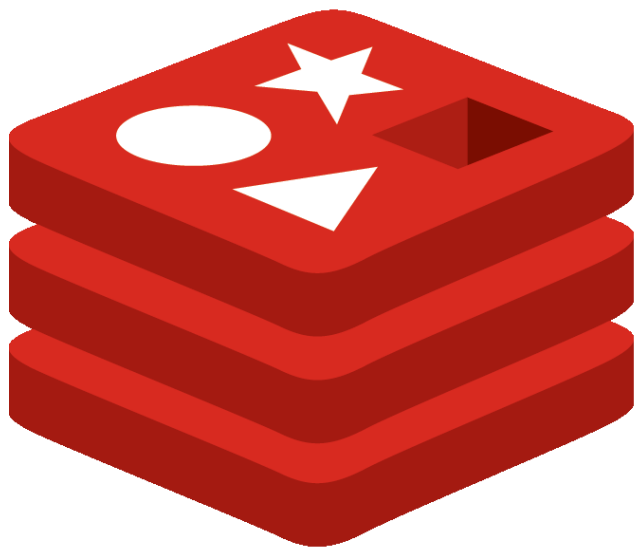
Strategy 3: Local Messages

- Why it is better in node:
 - In ~100 lines of JS...
 - Messages aren't lost when server dies
 - Webserver process unbound by email sending
 - Error handling, Throttling, Queuing and retries!
 - Offline support?
 - Why it is still a bad idea:
 - Bound to one host
-

Strategy 5: Remote Queues

- Observability
 - how long is the queue?
 - how long does an item wait in the queue?
 - ops stuff
 - Redundancy
 - Backups
 - Clustering
 - ops stuff
-

A Quick Aside



redis

REDIS IS REALLY AWESOME

Data Structures for a MVP Queue:

- Array
 - Push, Pop, Length

I guess that's it...

Data Structures for a good Queue:

- Array
 - Push, Pop, Length
 - Hash (key types: string, integer, hash)
 - Set, Get, Exists
 - Sorted Set
 - Exists, Add, Remove
-

Data Structures for a Good Queue

RESQUE ([node-resque](#))

```
var queue = new NR.queue({connection: connectionDetails}, jobs, function(){
  queue.enqueue('math', "add", [1,2]);
  queue.enqueue('math', "add", [2,3]);
  queue.enqueueIn(3000, 'math', "subtract", [2,1]);
});
```

Data Structures for a Good Queue



Refresh Commands More...

127.0.0.1:6379:0

resque:* (2)

queue:* (1)

emailQueue (1)

queues (1)

Add New Value...

Delete Key



0

Goto Index



Value

```
0 {"class":"sendEmail","queue":"emailQueue","args":
  [{"to":"evantahler@gmail.com","subject":"hello_from_node","text":"hello_again"}]}
```

Queues

The list below contains all the registered queues with the number of jobs currently in the queue. Select a queue from above to view all jobs currently pending on the queue.

Name	Jobs
<u>emailQueue</u>	1
<u>failed</u>	0

0 of 0 Workers Working

The list below contains all workers which are currently running a job.

	Where	Queue	Processing
Nothing is happening right now...			


```
var jobs = {  
  sendEmail: function(data, callback){  
    var email = {  
      from:    require('./emailUsername'),  
      to:      data.to,  
      subject: data.subject,  
      text:    data.text,  
    };  
  
    transporter.sendMail(email, function(error, info){  
      callback(error, {email: email, info: info});  
    });  
  }  
};
```

```
if(process.env.ROLE === 'web'){  
  var queue;
```

```
  var server = function(req, res){  
    var urlParts = req.url.split('/');  
    var email = {  
      to:      urlParts[1],  
      subject: urlParts[2],  
      text:    urlParts[3]}  
  };
```

```
  queue.enqueue('emailQueue', "sendEmail", email);
```


```
  var response = {email: email};  
  res.writeHead(200, {'Content-Type': 'application/json'});  
  res.end(JSON.stringify(response, null, 2));  
};
```

```
queue = new NR.queue({connection: connectionDetails}, jobs, function(){  
  http.createServer(server).listen(httpPort);  
  console.log('server running');  
});  
}
```

IPC!



**Connect before
server start**



```
if(process.env.ROLE === 'worker'){  
  var worker = new NR.worker({connection: connectionDetails, queues: ['emailQueue']}, jobs, function(){  
    worker.workerCleanup();  
    worker.start();  
  });
```

**Really
Simple**

```
worker.on('start', function(){ console.log("worker started"); });  
worker.on('end', function(){ console.log("worker ended"); });  
worker.on('cleaning_worker', function(worker, pid){ console.log("cleaning old worker " + worker); });  
worker.on('poll', function(queue){ console.log("worker polling " + queue); });  
worker.on('job', function(queue, job){ console.log("working job " + queue + " " + JSON.str); });  
worker.on('reEnqueue', function(queue, job, plugin){ console.log("reEnqueue job (" + plugin + "); "); });  
worker.on('success', function(queue, job, result){ console.log("job success " + queue + " " + JSON.str); });  
worker.on('failure', function(queue, job, failure){ console.log("job failure " + queue + " " + JSON.str); });  
worker.on('error', function(queue, job, error){ console.log("error " + queue + " " + JSON.str); });  
worker.on('pause', function(){ console.log("worker paused"); });
```

Tons of optional status events



So what is special about node.js here?

Queue Workers @ Node

- The event loops is great for processing **all** non-blocking events, not just web servers.
 - Most Background jobs are non-blocking events
 - Update the DB, Talk to this external service, etc
 - So node can handle **many** of these at once per process!
-

```
var multiWorker = new NR.multiWorker({
  connection: connectionDetails,
  queues: ['slowQueue'],
  minTaskProcessors: 1,
  maxTaskProcessors: 100,
  checkTimeout: 1000,
  maxEventLoopDelay: 10,
  toDisconnectProcessors: true,
}, jobs, function(){

  // normal worker emitters
  multiWorker.on('start', function(workerID) {
  multiWorker.on('end', function(workerID) {
  /// ...

  multiWorker.on('internalError', function(error){
  multiWorker.on('multiWorkerAction', function(verb, d

  multiWorker.start();
});
```

How can you tell the CPU is pegged?

// inspired by <https://github.com/tj/node-blocked>

```
module.exports = function(limit, interval, fn) {  
  var start = process.hrtime();
```

```
  setInterval(function(){  
    var delta = process.hrtime(start);  
    var nanosec = delta[0] * 1e9 + delta[1];  
    var ms = nanosec / 1e6;  
    var n = ms - interval;  
    if (n > limit){  
      fn(true, Math.round(n));  
    }else{  
      fn(false, Math.round(n));  
    }  
    start = process.hrtime();  
  }, interval).unref();  
};
```

+

process.setImmediate()

Example Time!

```
var blockingSleep = function(naptime){
  var sleeping = true;
  var now = new Date();
  var alarm;
  var startingMSeconds = now.getTime();
  while(sleeping){
    alarm = new Date();
    var alarmMSeconds = alarm.getTime();
    if(alarmMSeconds - startingMSeconds > naptime){ sleeping = false }
  }
}
```

```
var jobs = {
  "slowSleepJob": {
    plugins: [],
    pluginOptions: {},
    perform: function(callback){
      var start = new Date().getTime();
      setTimeout(function(){
        callback(null, (new Date().getTime() - start) );
      }, 1000);
    },
  },
  "slowCPUJob": {
    plugins: [],
    pluginOptions: {},
    perform: function(callback){
      var start = new Date().getTime();
      blockingSleep(1000);
      callback(null, (new Date().getTime() - start) );
    },
  },
};
```


Strategy 5: Remote Queues

- Why it is better in node:
 - In addition to persistent storage and multiple server/process support, you get CPU scaling and Throttling very simply!
 - Node also has tooling (domains) around async exceptions which other languages lack
 - Integrates well with the resque/sidekiq pattern
 - This might finally be a good idea!
-

THANKS!

- These Slides
 - goo.gl/yUuApo
 - Supporting Project:
 - https://github.com/evantahler/background_jobs_node
 - Node-Resque:
 - <https://github.com/taskrabbit/node-resque>
 - [MultiWorker Example](#)
-

Bonus Slides to follow

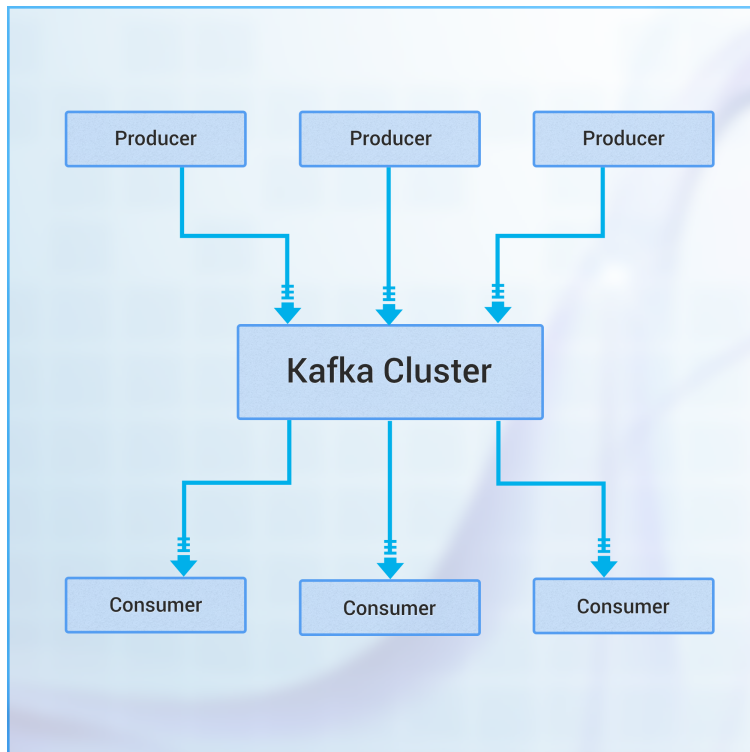
Strategy 4: Remote Messages

- Styles:
 - **Synchronous-processing**
 - Can provide messaging to the client about success
 - But the client still has to wait...
 - **Asynchronous-processing**
 - Just like our cluster example, but now we can separate **servers** and not just **processes**
-

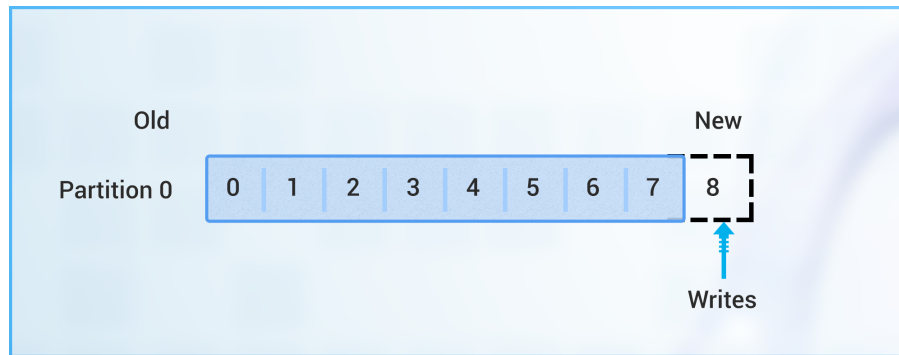
Strategy 4: Remote Messages

- **Synchronous-processing** doesn't seem help too much (unless there are *OPS* considerations)
 - How can we build a persistent **Asynchronous-processing** app?
 - We'll need that app to respond with status
 - Job Started, job failed, job succeeded...
 - We'll use a **Remote Queue!**
-

Strategy 6: Event Bus



(Watch or Poll) vs Push



<http://blog.qburst.com/2014/06/apache-kafka/>